



A1: Text Mining Report on Airbnb

**Uncovering Market Insights for Hosts and Platform
Optimization**

Mar 13, 2025

PRESENTED TO

PRESENTED BY

Professor Thomas Kurnicki

Tirthankar Mukherjee

Business Analysis with Unstructured Data - DAT-7471- FMBANDD1

Hult International Business School

Executive Summary

This comprehensive analysis of Airbnb listing data highlights several key factors significantly influencing booking rates and pricing strategies. Listings with neighborhood descriptions that exude positive sentiment, particularly those featuring phrases like “walking distance,” “metro station,” and “restaurants bars,” tend to attract more bookings. Pricing trends indicate that entire homes and apartments, primarily ‘serviced apartments and villas’ priced above \$600, command premium rates. In contrast, shared rooms are priced considerably lower. Property size metrics, such as the number of bedrooms and bathrooms, strongly correlate with price points. While host response rates correlate negatively with pricing, they are likely crucial for enhancing guest satisfaction.

Based on these insights, Airbnb management is recommended to implement the following four strategic initiatives:

Optimize Pricing Strategies: Encourage hosts to focus on premium property types and utilize data-driven approaches to set competitive cleaning fees.

Enhance Neighborhood Descriptions: Guide hosts to incorporate high-impact phrases related to accessibility and amenities to attract more bookings.

Improve Host Performance: Develop training programs and offer superhost incentives to elevate host responsiveness and service quality.

Leverage Text Analytics: Use advanced text analytics to provide personalized guest recommendations and improve service offerings.

These strategies will enable hosts to maximize their revenue potential while aligning their listings closely with guest preferences for convenience, accessibility, and dining options.

Visualizations & Text Mining Framework with Key Insights

1. “Average Price by Room Type and Property Type” visualization: This bar chart lets viewers quickly identify which room type and property type combinations command the highest prices on the Airbnb platform. Entire home listings have higher average prices than private or shared rooms. At the same time, Serviced Apartments and Villas are the most expensive property types under "Private Room" based on the highest bars on the chart, followed by hotels and apartments. (*Appendix 1*)

2. "Sentiment Analysis of Airbnb Summaries" is a series of bar charts, each representing the sentiment score derived from Airbnb summaries using different sentiment analysis methods. The AFINN lexicon in red has a sentiment score of approximately 15,000, identifying a relatively high positive emotional tone in the Airbnb summaries. The Bing lexicon in green and the NRC lexicon in blue, showing a sentiment score of around 12,000 and approximately 10,000, indicate a positive overall sentiment. (*Appendix 2*)

3. “Correlation Heatmap of Numeric Variables” visualization: This correlation heatmap displays the pairwise correlations between different numeric variables in the Airbnb dataset. Strong positive correlations are visible between variables related to property size, such as bathrooms and beds. This indicates that listings with more bedrooms tend to have more bathrooms and can accommodate more guests. (*Appendix 3*)

4. “Superhost vs. Non-Superhost Performance” visualization: The bar chart highlights that Superhosts positively impact guest engagement, as indicated by the higher average number of reviews and higher average review scores. However, the observation that the superhost has a smaller average price provides potential business insight implications: to promote superhosts and incentivize more hosts to go above and beyond. (*Appendix 4*)

5. “Most Common Words in Airbnb Descriptions” Word Cloud: The word cloud summarizes key themes used in Airbnb listing descriptions, providing valuable insights for hosts to optimize their

listings and attract more bookings like prominence of location, accessibility, and essential amenities indicating customer preference when searching for accommodations. (Appendix 5)

6. Bigram Network of Neighborhood Descriptions: The bigram network visualization effectively summarizes frequent relationships in the text data of neighborhood descriptions. It reveals clusters of terms like “coffee shops” and "Bondi Beach," offering insights into what makes neighborhoods desirable in Airbnb listings. The visualization can inform hosts about what features of their location they should emphasize in their descriptions to attract more bookings. (Appendix 6)

7. “Top 20 Bigrams in Neighborhood Descriptions” visualization: This horizontal bar chart highlights the importance of walkability, proximity to key amenities, and locations in describing Airbnb neighborhoods and can influence how listings are worded. The bigram "walking distance" and “5 minutes” are used the most, clearly standing out as key elements in neighborhood descriptions. (Appendix 7)

Airbnb Data Analysis Dashboard Summary

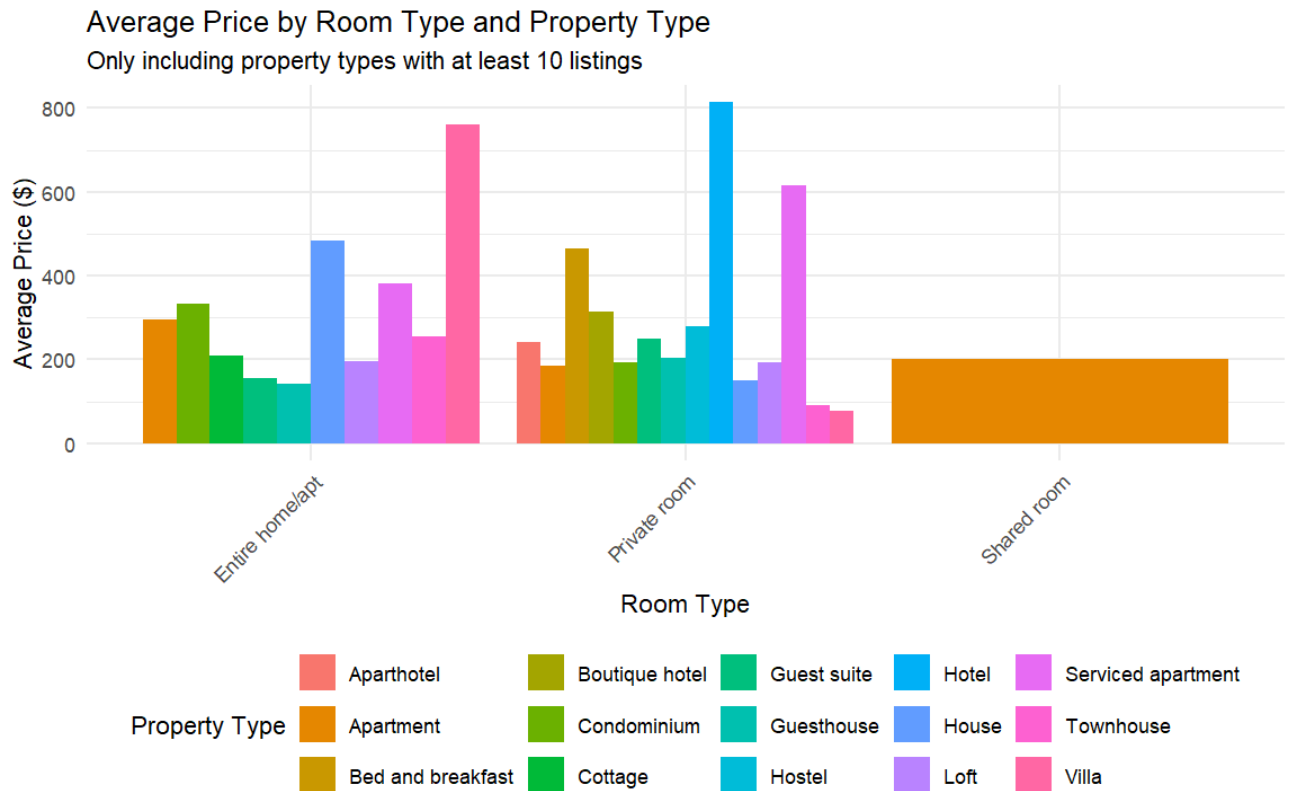
The dashboard combines filtering capabilities with visual and tabular data presentations to help users understand pricing patterns, review trends, and content themes data through multiple analytical perspectives. The interface is divided into two main components:

1. Sidebar Controls- This section includes a *room type selector*, *price range slider*, and *superhosts-only* checkbox.
2. Main Content Area- This section uses a tabbed interface to organize different analytical views like a price analysis tab (to compare average prices of rooms and properties), reviews tab (visualize review counts by room type), text analysis tab (a word cloud visualizing the most frequent terms in listing descriptions), and data table tab (comprehensive listing details including name, type, price, and amenities). (Appendix 8.A, 8.B)

This dashboard allows managers to make strategic decisions about Airbnb market positioning.

Appendix

Appendix 1: Description of the "Average Price by Room Type and Property Type" visualization



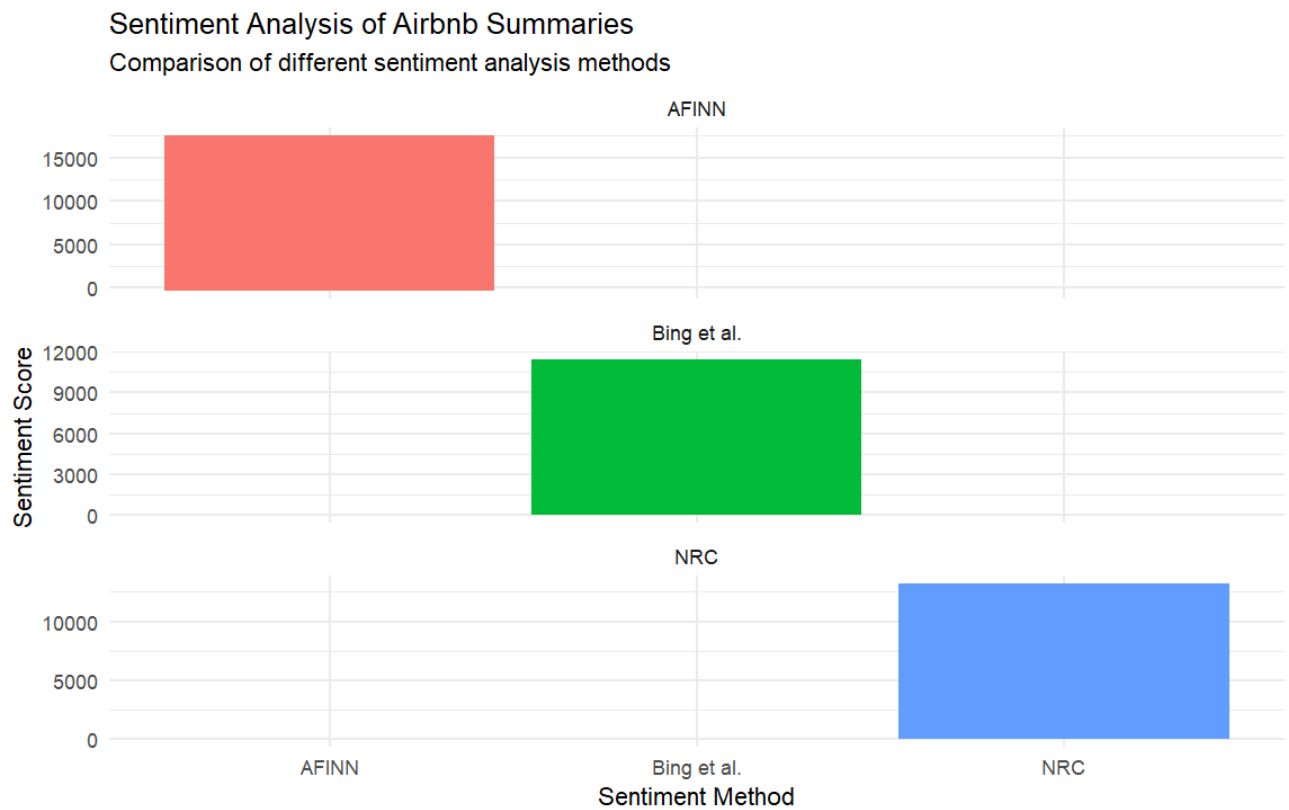
Axes:

- **X-axis:** Represents Room Type. The three room types are "Entire home/apt," "Private room," and "Shared room."
- **Y-axis:** Represents Average Price (\$). The scale ranges from 0 to 800.

Data Representation:

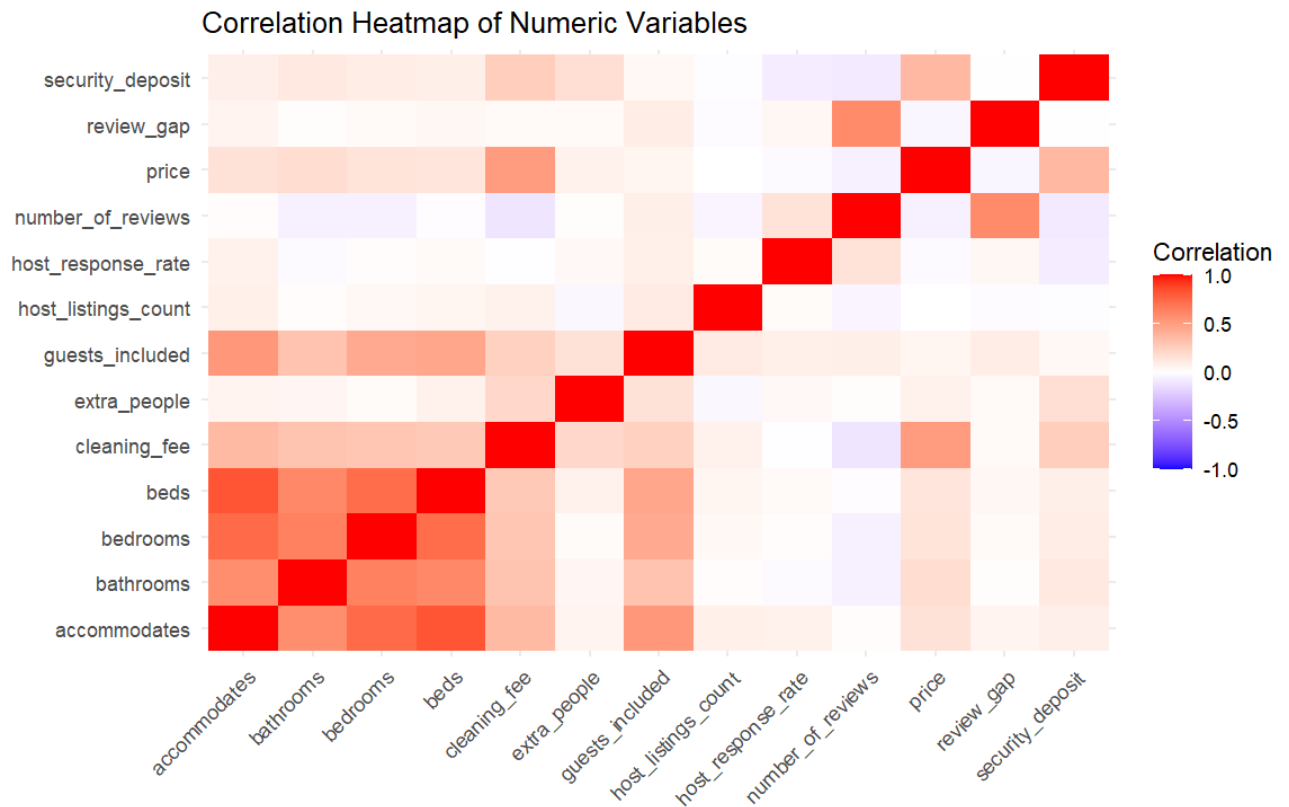
- Each bar represents a specific combination of room type and property type.
- The height of each bar corresponds to the average price for that combination.
- Bars are color-coded to distinguish between different property types, with a legend provided beneath the chart.

Appendix 2: Description of the sentiment analysis visualization for the Airbnb dataset:



The y-axis represents the **Sentiment Score**, and the x-axis represents the **Sentiment Method**.

Appendix 3: Description of the "Correlation Heatmap of Numeric Variables" visualization for the Airbnb dataset:



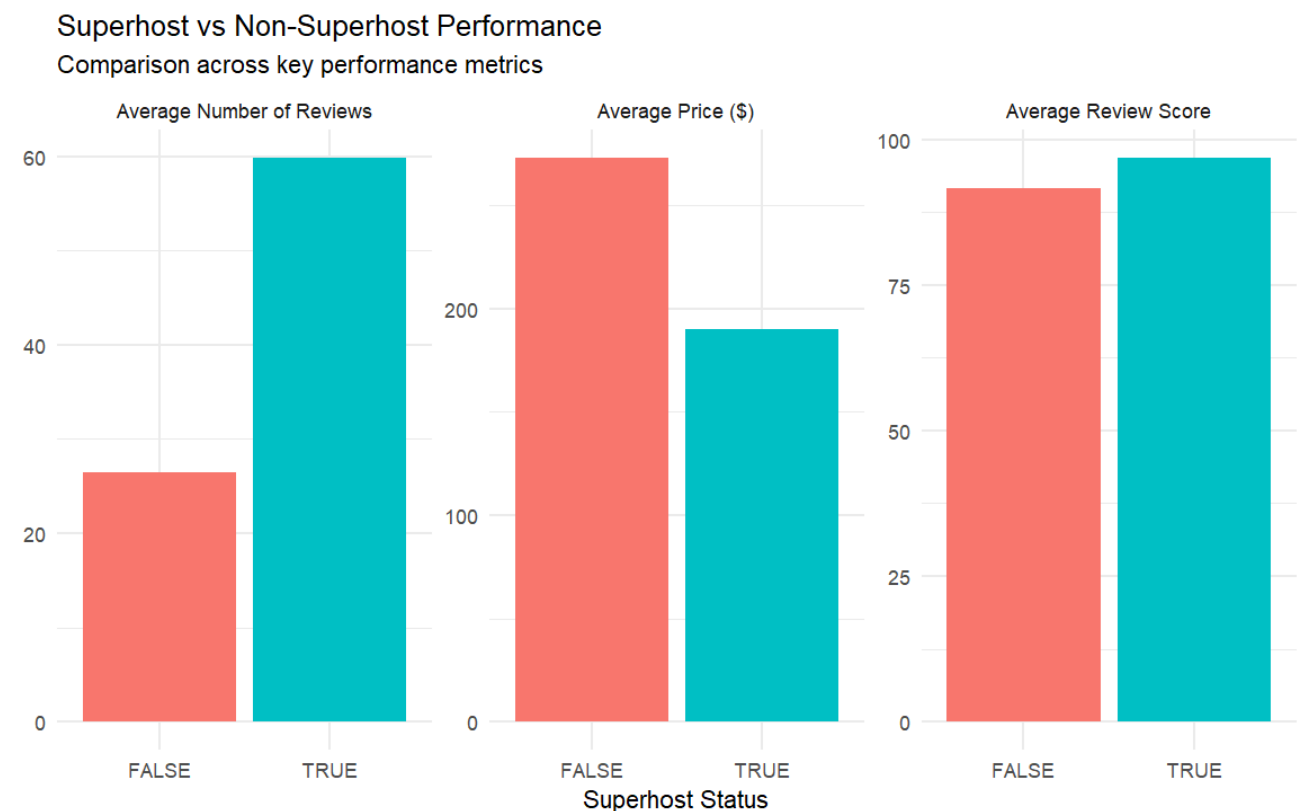
Cells: Each cell at the intersection of a row and a column represents the correlation coefficient between the two corresponding variables.

Color Coding: The cells are colored according to the strength and direction of the correlation:

- **Red** indicates a positive correlation. The intensity of the red color represents the strength of the positive correlation (darker red = stronger positive correlation).
- **White:** Indicates a correlation close to zero (no or very weak correlation).
- **Blue** Indicates a negative correlation. The intensity of the blue color represents the strength of the negative correlation (darker blue = stronger negative correlation).

Color Scale: A color scale (legend) is provided to the right of the heatmap, mapping the color intensities to correlation values ranging from -1.0 to 1.0.

Appendix 4: “Superhost vs Non-Superhost Performance: Comparison across key performance metrics”, comparing the performance of Airbnb listings based on whether the host is a Superhost (TRUE) or not (FALSE)



X-axis: Superhost Status (FALSE/TRUE).

Average Number of Reviews: **Y-axis:** *Average Number of Reviews.* The turquoise bar is significantly higher than the coral bar, indicating that listings managed by Superhosts tend to have a much higher average number of reviews than those managed by non-Superhosts.

Average Price: **Y-axis:** *Average Price.* On average, the listings managed by non-Superhosts have a higher price than those managed by Superhosts, possibly to break even owing to fewer customers.

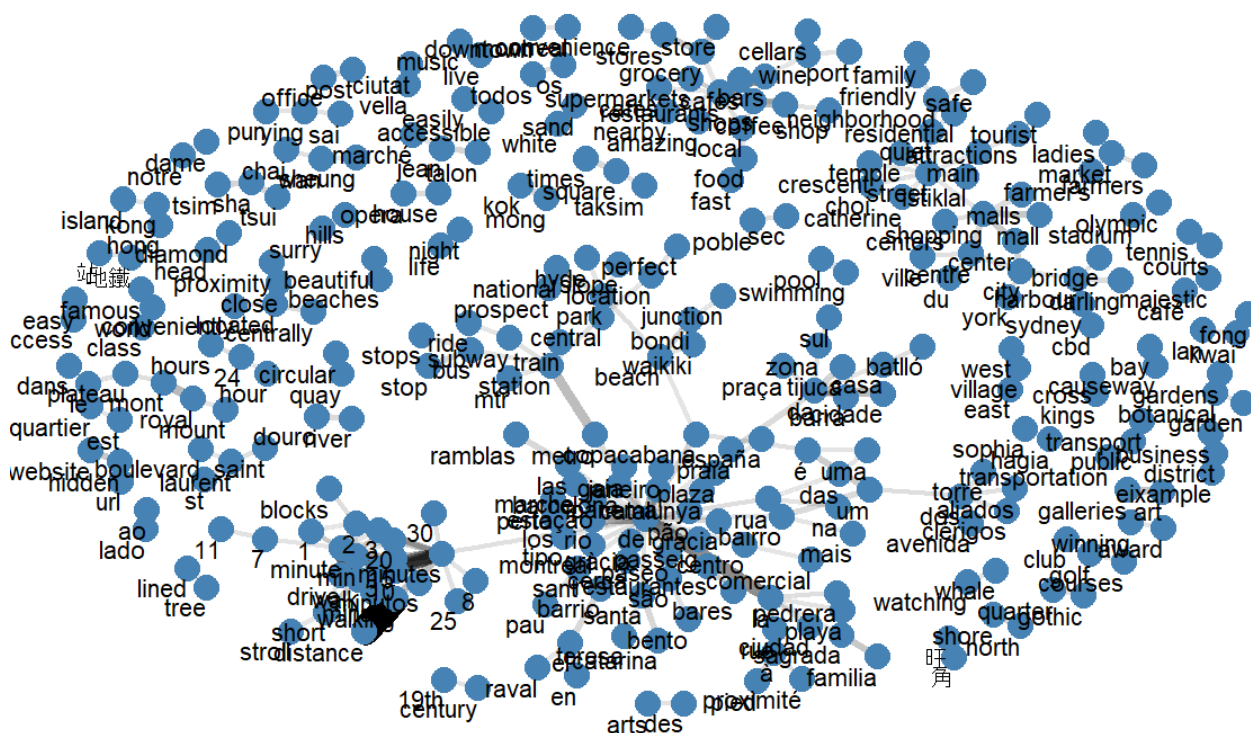
Average Review Score: **Y-axis:** *Average Review Score.* The bar graph suggests that listings managed by Superhosts generally have higher average review scores.



Dominant Words: Due to its size, the most prominent word is “apartment.” Other significant words include “minutes,” “located,” “restaurants,” “kitchen,” and “bedroom.” These suggest that descriptions often focus on property type, location, amenities, and the number of bedrooms.

Descriptive Adjectives: Adjectives like “beautiful,” “cozy,” “comfortable,” “modern,” and “quiet” give insight into how hosts are describing their properties to attract customers.

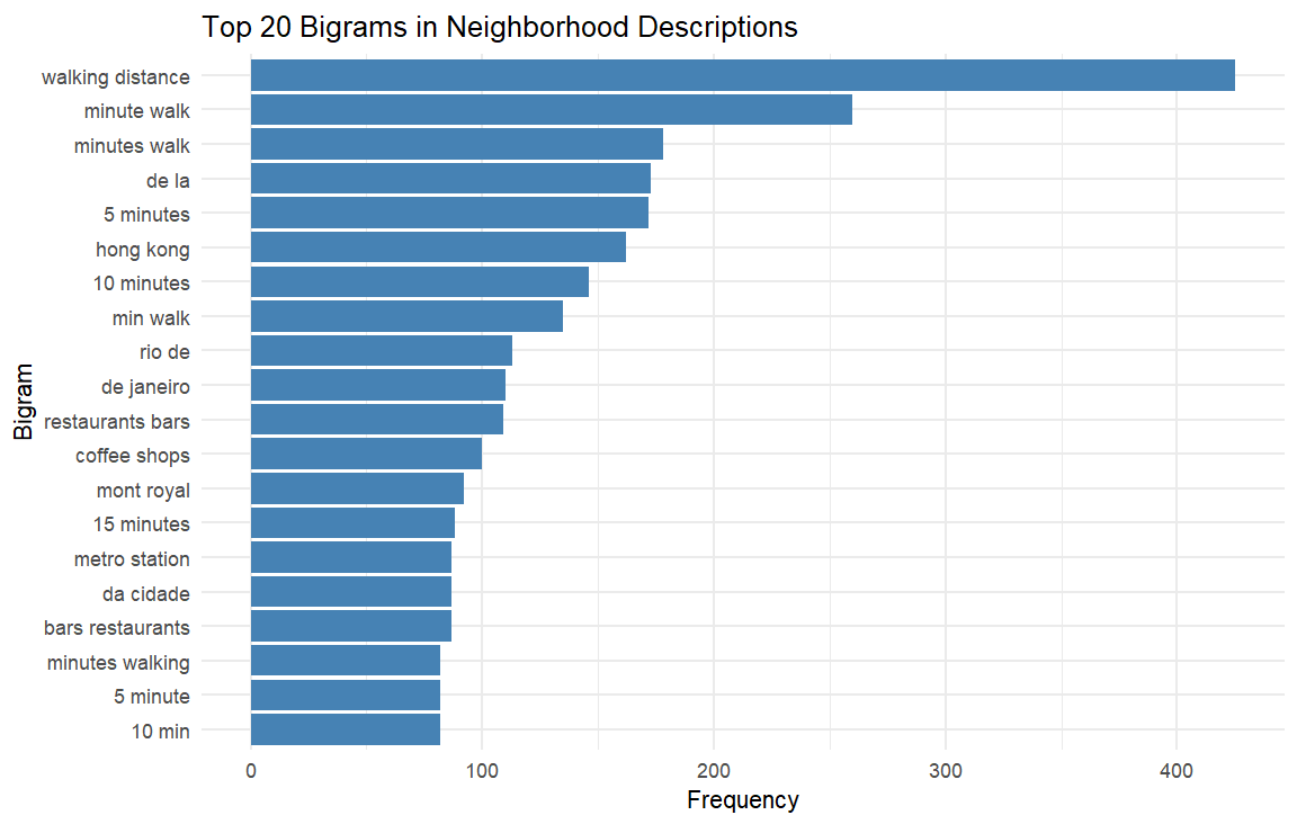
Showing bigrams appearing at least 15 times



Edges: Lines connect the nodes, representing the co-occurrence of these words or phrases within the neighborhood descriptions. The thickness of the line suggests the frequency of their co-occurrence, with thicker lines indicating more frequent pairings.

Layout: The nodes are arranged in a force-directed layout where connected nodes are drawn closer together and disconnected nodes are pushed further apart. This allows clusters of related terms to emerge.

Appendix 7: Horizontal bar chart titled “Top 20 Bigrams in Neighborhood Descriptions”.



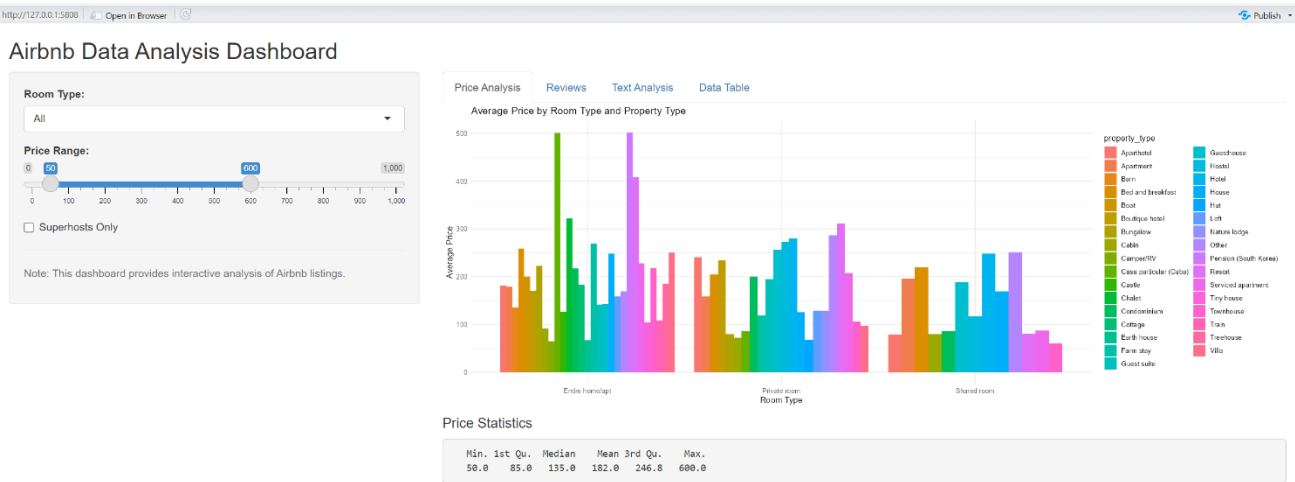
Y-axis: The y-axis lists the 20 most frequent bigrams. Each bigram is a pair of words often appearing together in the neighborhood descriptions. Examples include "walking distance," "minute walk," "restaurants bars," and location-specific phrases like "hong kong," "rio de," and "de janeiro." The bigrams are ordered from top to bottom according to frequency (highest to lowest).

X-axis: The x-axis represents the frequency or count of each bigram in the dataset. The scale extends from 0 to 400.

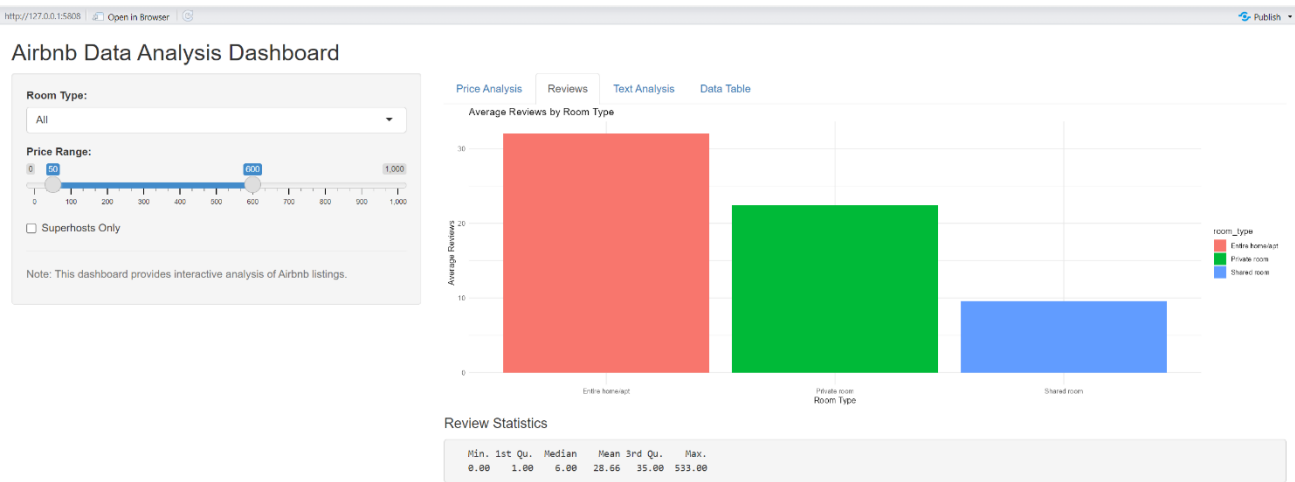
Bars: Each bigram has a horizontal bar extending from the y-axis to a point corresponding to its frequency on the x-axis. The bars are all colored in blue.

Appendix 8.A: Description of the Airbnb Data Analysis Dashboard

Price Analysis Tab:

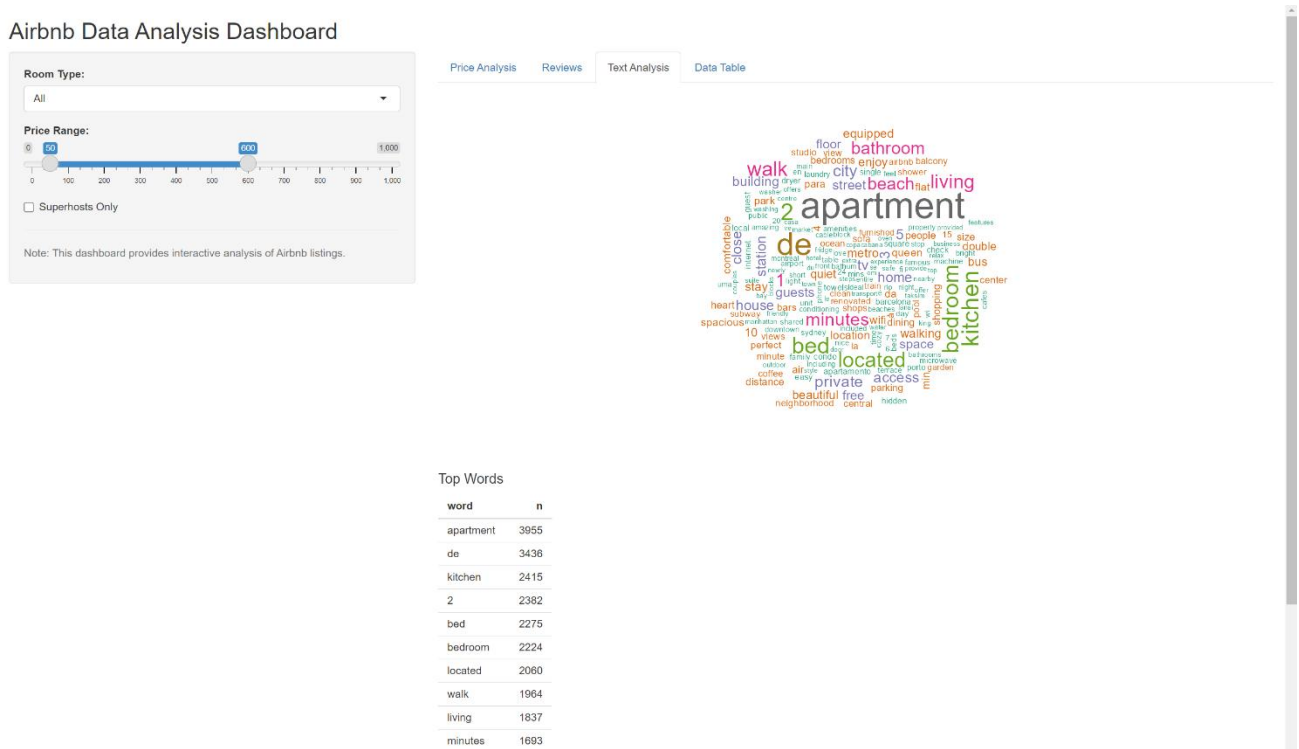


Reviews Tab:

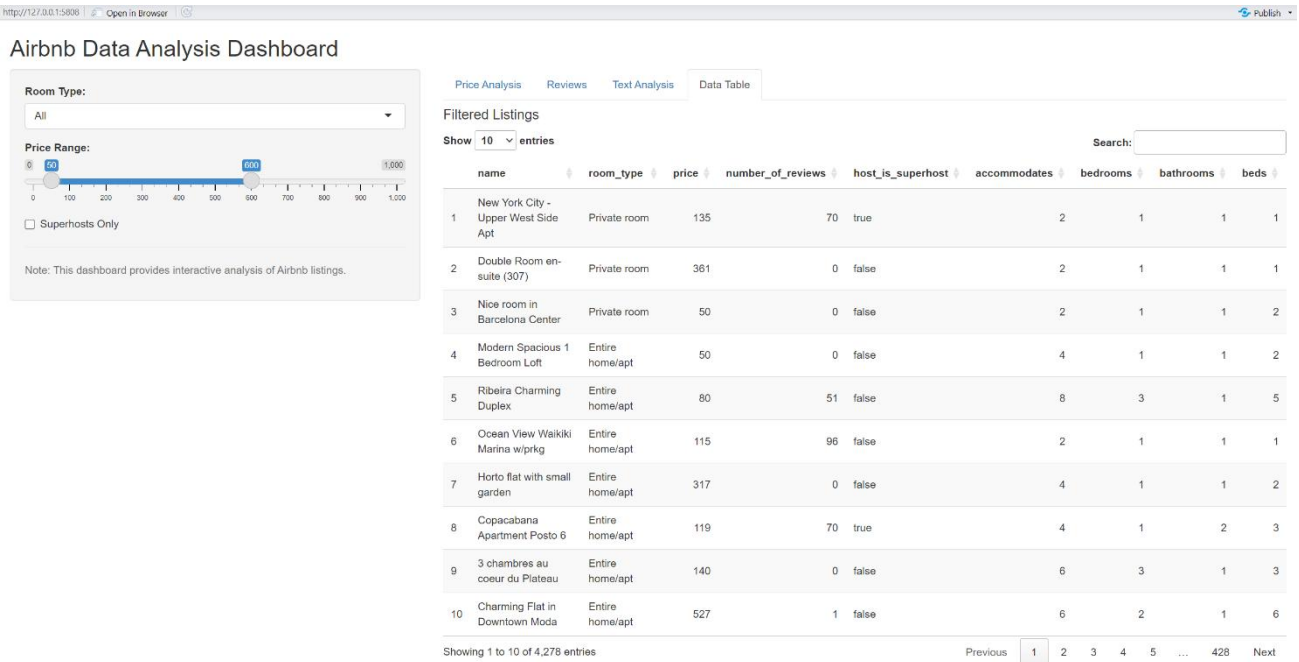


Appendix 8.B: Description of the Airbnb Data Analysis Dashboard

Text Analysis Tab:

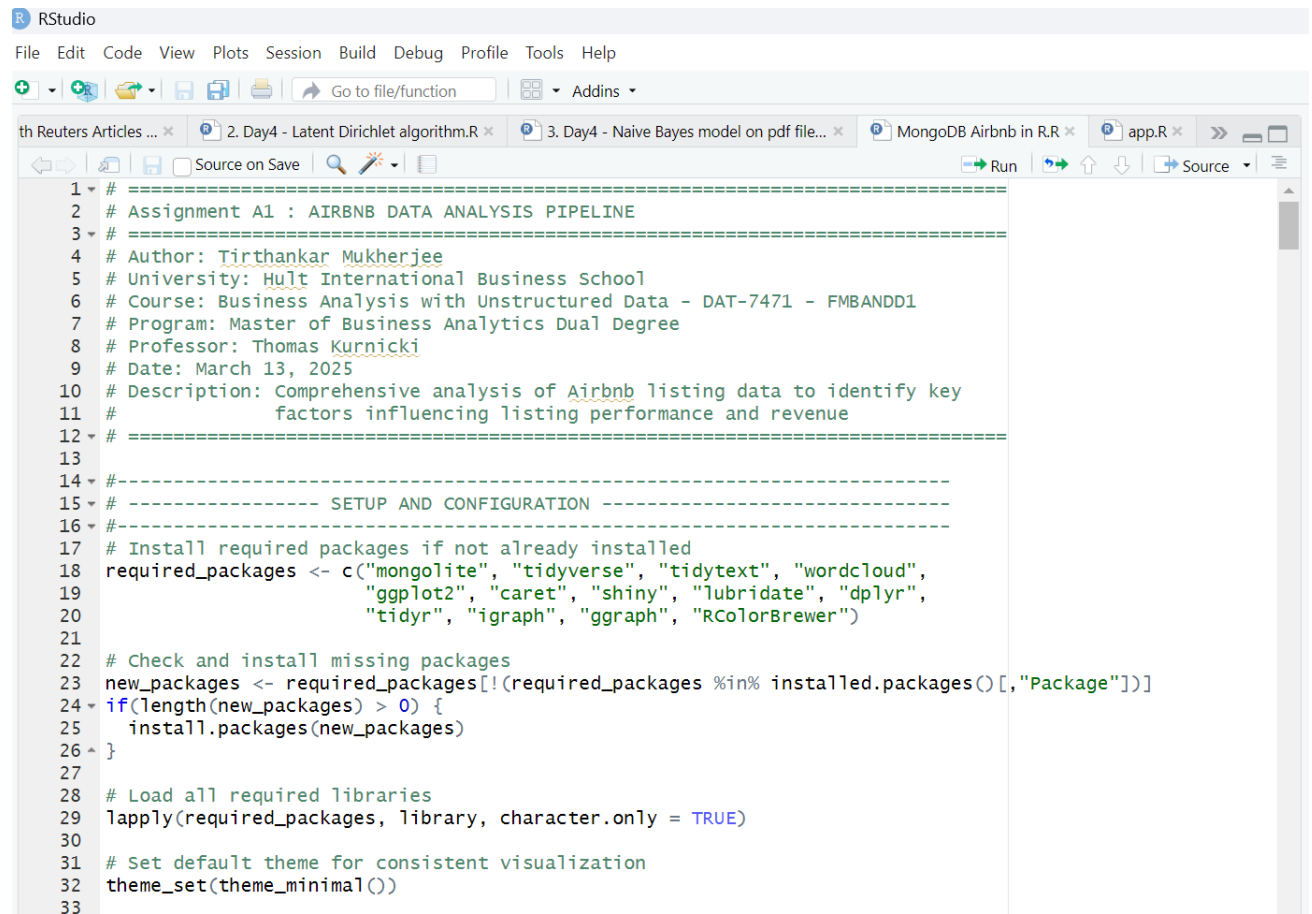


Data Table Tab:



Appendix 9: R-Script

Appendix 9.A: Setup and Configuration



The screenshot shows the RStudio interface with a script editor containing R code for setting up the environment. The code includes comments about the assignment and author, followed by a section for installing and loading required packages.

```
1 # =====
2 # Assignment A1 : AIRBNB DATA ANALYSIS PIPELINE
3 # =====
4 # Author: Tirthankar Mukherjee
5 # University: Hult International Business School
6 # Course: Business Analysis with Unstructured Data - DAT-7471 - FMBANDD1
7 # Program: Master of Business Analytics Dual Degree
8 # Professor: Thomas Kurnicki
9 # Date: March 13, 2025
10 # Description: Comprehensive analysis of Airbnb listing data to identify key
11 #               factors influencing listing performance and revenue
12 # =====
13
14 # -----
15 # ----- SETUP AND CONFIGURATION -----
16 # -----
17 # Install required packages if not already installed
18 required_packages <- c("mongolite", "tidyverse", "tidytext", "wordcloud",
19                       "ggplot2", "caret", "shiny", "lubridate", "dplyr",
20                       "tidyr", "igraph", "ggraph", "RColorBrewer")
21
22 # Check and install missing packages
23 new_packages <- required_packages[!(required_packages %in% installed.packages()[,"Package"])]
24 if(length(new_packages) > 0) {
25   install.packages(new_packages)
26 }
27
28 # Load all required libraries
29 lapply(required_packages, library, character.only = TRUE)
30
31 # Set default theme for consistent visualization
32 theme_set(theme_minimal())
33
```

Appendix 9.B: Data Connection and Loading

```
#-----  
# ----- DATA CONNECTION AND LOADING -----  
#-----  
# MongoDB connection details  
# Note: In production, consider storing credentials in environment variables  
connection_string <- 'mongodb+srv://mukherjeetirthankar:tirthA25@a1textmining.9gdz0.mongodb.net/?retryw  
  
# Initialize MongoDB connection  
airbnb_collection <- mongo(  
  collection = "listingsAndReviews",  
  db = "sample_airbnb",  
  url = connection_string  
)  
  
# Function to clean and preprocess Airbnb data  
clean_airbnb_data <- function(data) {  
  # Data cleaning and type conversion  
  data %>%  
    mutate(  
      # Convert currency string to numeric  
      price = as.numeric(gsub("[\\$,]", "", price)),  
  
      # Convert date strings to Date objects  
      last_review = as.Date(last_review),  
      first_review = as.Date(first_review),  
  
      # Calculate duration between first and last review in days  
      review_gap = as.numeric(difftime(last_review, first_review, units = "days")),  
  
      # Extract and clean host response rate  
      host_response_rate = as.numeric(gsub("%", "", host$host_response_rate)) / 100,  
  
      # Convert listing details to numeric  
      accommodates = as.numeric(accommodates),  
      bedrooms = as.numeric(bedrooms),  
      bathrooms = as.numeric(bathrooms),  
      beds = as.numeric(beds),  
  
      # Extract additional host information  
      host_is_superhost = host$host_is_superhost,  
      host_identity_verified = host$host_identity_verified,  
      host_listings_count = as.numeric(host$host_listings_count)  
    )  
}  
  
# Fetch and preprocess all data  
# Warning: This operation may require significant memory  
airbnb <- airbnb_collection$find() %>% clean_airbnb_data()
```

Appendix 9.C: Exploratory Data Analysis

```
87 ▾ #-----  
88 ▾ #----- EXPLORATORY DATA ANALYSIS -----  
89 ▾ #-----  
90  
91 # Function to summarize dataset  
92 ▾ summarize_dataset <- function(data) {  
93   # Basic dataset summary  
94   cat("Dataset Summary:\n")  
95   cat("Number of listings:", nrow(data), "\n")  
96   cat("Number of variables:", ncol(data), "\n")  
97   cat("Date range:", min(data$first_review, na.rm = TRUE), "to",  
98     max(data$last_review, na.rm = TRUE), "\n")  
99  
100  # Missing values summary  
101  missing_values <- data %>%  
102    summarise(across(everything(), ~sum(is.na(.)))) %>%  
103    pivot_longer(everything(), names_to = "variable", values_to = "missing_count") %>%  
104    filter(missing_count > 0) %>%  
105    arrange(desc(missing_count))  
106  
107 ▾  if(nrow(missing_values) > 0) {  
108    cat("\nVariables with missing values:\n")  
109    print(missing_values)  
110 ▴  }  
111 ▴ }  
112  
113 # Run initial data summary  
114 summarize_dataset(airbnb)  
115  
116
```


Appendix 9.D: Price Analysis

```
117 #-----
118 # ----- PRICE ANALYSIS -----
119 #-----
120 # Analyze how price varies by room type and property type
121 price_analysis <- airbnb %>%
122   # Group data for aggregation
123   group_by(room_type, property_type) %>%
124   # Calculate metrics for each group
125   summarise(
126     avg_price = mean(price, na.rm = TRUE),
127     median_price = median(price, na.rm = TRUE),
128     min_price = min(price, na.rm = TRUE),
129     max_price = max(price, na.rm = TRUE),
130     total_listings = n(),
131     avg_reviews = mean(number_of_reviews, na.rm = TRUE),
132     .groups = "drop" # Drop grouping after summarization
133   ) %>%
134   # Filter out uncommon property types for clearer visualization
135   filter(total_listings >= 10)
136
137 # Visualization: Average Price by Room Type and Property Type
138 plot_price_by_room_property <- function(data) {
139   ggplot(data, aes(x = room_type, y = avg_price, fill = property_type)) +
140     geom_bar(stat = "identity", position = "dodge") +
141     labs(
142       title = "Average Price by Room Type and Property Type",
143       subtitle = "Only including property types with at least 10 listings",
144       x = "Room Type",
145       y = "Average Price ($)",
146       fill = "Property Type"
147     ) +
148     theme(
149       axis.text.x = element_text(angle = 45, hjust = 1),
150       legend.position = "bottom",
151       legend.box = "horizontal"
152     )
153 }
154
155 # Generate and display the price analysis plot
156 plot_price_by_room_property(price_analysis)
157
158
```

Appendix 9.E: Host Performance Analysis

```
159 #-----
160 #----- HOST PERFORMANCE ANALYSIS -----
161 #-----
162 # Analyze how host characteristics affect listing performance
163 host_performance <- airbnb %>%
164   # Filter out entries with missing response rate
165   filter(!is.na(host_response_rate)) %>%
166   # Group by superhost status
167   group_by(host_is_superhost) %>%
168   # Calculate metrics for each group
169   summarise(
170     avg_response_rate = mean(host_response_rate, na.rm = TRUE),
171     avg_price = mean(price, na.rm = TRUE),
172     avg_reviews = mean(number_of_reviews, na.rm = TRUE),
173     avg_review_scores = mean(review_scores$review_scores_rating, na.rm = TRUE),
174     total_listings = n(),
175     .groups = "drop"
176   )
177
178 # Visualization: Superhost vs Non-Superhost Performance
179 plot_superhost_performance <- function(data) {
180   # Prepare data for visualization by pivoting to long format
181   plot_data <- data %>%
182     select(host_is_superhost, avg_price, avg_reviews, avg_review_scores) %>%
183     pivot_longer(
184       cols = c(avg_price, avg_reviews, avg_review_scores),
185       names_to = "metric",
186       values_to = "value"
187     ) %>%
188     mutate(
189       # Create readable labels for the metrics
190       metric = case_when(
191         metric == "avg_price" ~ "Average Price ($)",
192         metric == "avg_reviews" ~ "Average Number of Reviews",
193         metric == "avg_review_scores" ~ "Average Review Score",
194         TRUE ~ metric
195       )
196     )
197
198   # Create the plot
199   ggplot(plot_data, aes(x = host_is_superhost, y = value, fill = host_is_superhost)) +
200     geom_bar(stat = "identity") +
201     facet_wrap(~ metric, scales = "free_y") +
202     labs(
203       title = "Superhost vs Non-Superhost Performance",
204       subtitle = "Comparison across key performance metrics",
205       x = "Superhost Status",
206       y = NULL
207     ) +
208     theme(legend.position = "none")
209 }
210
211 # Generate and display the host performance plot
212 plot_superhost_performance(host_performance)
```

Appendix 9.F: Text Mining on Descriptions

```
215 #-----
216 #----- TEXT MINING ON DESCRIPTIONS -----
217 #-----
218 # Process and analyze listing descriptions to identify key terms
219
220 # Function to process text and remove stop words
221 process_text <- function(data, text_column, id_column = "name") {
222   data %>%
223     select(!sym(id_column), !sym(text_column)) %>%
224     # Filter out missing text values
225     filter(!is.na(!sym(text_column))) %>%
226     # Tokenize text into individual words
227     unnest_tokens(word, !sym(text_column)) %>%
228     # Remove stop words that don't add meaning
229     anti_join(stop_words) %>%
230     # Count word frequency
231     count(word, sort = TRUE)
232 }
233
234 # Process descriptions for analysis
235 description_tokens <- process_text(airbnb, "description")
236
237 # Create word cloud of most common words in descriptions
238 create_wordcloud <- function(token_data, max_words = 200, title = "Word Cloud") {
239   # Set up plotting device
240   par(mar = c(0, 0, 2, 0))
241
242   # Generate word cloud
243   wordcloud(
244     words = token_data$word,
245     freq = token_data$n,
246     max.words = max_words,
247     colors = brewer.pal(8, "Dark2"),
248     random.order = FALSE,
249     rot.per = 0.35,
250     scale = c(3, 0.5)
251   )
252
253   # Add title
254   title(main = title)
255 }
256
257 # Generate and display the description word cloud
258 create_wordcloud(
259   description_tokens,
260   max_words = 200,
261   title = "Most Common Words in Airbnb Descriptions"
262 )
263
```

Appendix 9.G: Sentiment Analysis

```
265 #----- SENTIMENT ANALYSIS ON LISTING SUMMARIES -----
266 #-----
267 # Analyze the sentiment of listing summaries to identify emotional tone
268
269 # Function to perform sentiment analysis using multiple lexicons
270 analyze_sentiment <- function(token_data) {
271   # AFINN sentiment analysis (numerical scores)
272   afinn_sentiment <- token_data %>%
273     inner_join(get_sentiments("afinn")) %>%
274     group_by(name) %>%
275     summarise(sentiment = sum(value), .groups = "drop") %>%
276     mutate(method = "AFINN")
277
278   # Bing and NRC sentiment analysis (categorical)
279   bing_and_nrc_sentiment <- bind_rows(
280     # Bing lexicon (positive/negative)
281     token_data %>%
282       inner_join(get_sentiments("bing")) %>%
283       mutate(method = "Bing et al."),
284
285     # NRC lexicon (positive/negative only)
286     token_data %>%
287       inner_join(get_sentiments("nrc")) %>%
288       filter(sentiment %in% c("positive", "negative")) %>%
289       mutate(method = "NRC")
290   ) %>%
291   # Count positive and negative words
292   count(method, sentiment) %>%
293   # Reshape data for visualization
294   pivot_wider(names_from = sentiment, values_from = n, values_fill = 0) %>%
295   # Calculate net sentiment
296   mutate(sentiment = positive - negative)
297
298   # Combine results from different methods
299   bind_rows(afinn_sentiment, bing_and_nrc_sentiment)
300 }
301
302 # Process and tokenize summaries
303 summary_tokens <- airbnb %>%
304   select(name, summary) %>%
305   filter(!is.na(summary)) %>%
306   unnest_tokens(word, summary) %>%
307   filter(!is.na(word)) %>%
308   anti_join(stop_words)
309
310 # Perform sentiment analysis
311 combined_sentiments <- analyze_sentiment(summary_tokens)
312
313 # Visualization: Sentiment Analysis Results
314 plot_sentiment_analysis <- function(sentiment_data) {
315   ggplot(sentiment_data, aes(method, sentiment, fill = method)) +
316     geom_col(show.legend = FALSE) +
317     facet_wrap(~method, ncol = 1, scales = "free_y") +
318     labs(
319
```

517:1 (Untitled) ↕

R Script ↕

Appendix 9.H: Neighborhood Bigram Analysis

```
331 #----- NEIGHBORHOOD BIGRAM ANALYSIS -----
332 #-----
333 # Analyze neighborhood descriptions using bigrams to identify location features
334
335 # Function to process text into n-grams
336 process_ngrams <- function(data, text_column, n = 2, id_column = "name", min_count = 1) {
337   # Create ngrams from text
338   ngrams <- data %>%
339     filter(!is.na(!!sym(text_column))) %>%
340     unnest_tokens(
341       ngram,
342       !!sym(text_column),
343       token = "ngrams",
344       n = n
345     )
346
347   # For bigrams, separate into component words and filter stop words
348   if (n == 2) {
349     ngrams <- ngrams %>%
350       separate(ngram, c("word1", "word2"), sep = " ") %>%
351       filter(!word1 %in% stop_words$word) %>%
352       filter(!word2 %in% stop_words$word) %>%
353       filter(!is.na(word1) & !is.na(word2)) %>%
354       count(word1, word2, sort = TRUE) %>%
355       filter(n >= min_count)
356   } else {
357     # For other n-grams, just count occurrences
358     ngrams <- ngrams %>%
359     count(ngram, sort = TRUE) %>%
360     filter(n >= min_count)
361   }
362
363   return(ngrams)
364 }
365
366 # Process neighborhood descriptions into bigrams
367 neighborhood_bigrams <- process_ngrams(
368   airtbnb,
369   "neighborhood_overview",
370   n = 2,
371   min_count = 10
372 )
373
374 # Visualization: Top Bigrams Bar Chart
375 plot_top_bigrams <- function(bigram_data, top_n = 20) {
376   bigram_data %>%
377     head(top_n) %>%
378     unite(bigram, word1, word2, sep = " ") %>%
379     ggplot(aes(x = reorder(bigram, n), y = n)) +
380     geom_col(fill = "steelblue") +
381     coord_flip() +
382     labs(
383       title = paste("Top", top_n, "Bigrams in Neighborhood Descriptions"),
384       x = "Bigram",
```

```
392 # Visualization: Bigram Network Graph
393 plot_bigram_network <- function(bigram_data, min_count = 10) {
394   # Filter bigrams by minimum count
395   filtered_bigrams <- bigram_data %>%
396     filter(n >= min_count)
397
398   # Create network graph
399   bigram_graph <- filtered_bigrams %>%
400     graph_from_data_frame()
401
402   # Plot the graph
403   ggraph(bigram_graph, layout = "fr") +
404     geom_edge_link(aes(edge_alpha = n, edge_width = n), show.legend = FALSE) +
405     geom_node_point(color = "steelblue", size = 5) +
406     geom_node_text(aes(label = name), vjust = 1.5, hjust = 1) +
407     labs(
408       title = "Bigram Network of Neighborhood Descriptions",
409       subtitle = paste("Showing bigrams appearing at least", min_count, "times")
410     ) +
411     theme_void()
412 }
413
414 # Generate and display the bigram network
415 plot_bigram_network(neighborhood_bigrams, min_count = 15)
416
```

Appendix 9.I: TF-IDF Analysis

```
418 #-----
419 #----- TF-IDF ANALYSIS -----
420 #-----
421
422 # Identify distinctive terms in neighborhood descriptions using TF-IDF
423
424 # Function to perform TF-IDF analysis on bigrams
425 analyze_tfidf <- function(data, text_column, grouping_var = "neighborhood_cleansed") {
426   # Process text into bigrams
427   bigrams <- data %>%
428     filter(!is.na(!sym(text_column)) & !is.na(!sym(grouping_var))) %>%
429     unnest_tokens(
430       bigram,
431       !!sym(text_column),
432       token = "ngrams",
433       n = 2
434     ) %>%
435     separate(bigram, c("word1", "word2"), sep = " ") %>%
436     filter(!word1 %in% stop_words$word) %>%
437     filter(!word2 %in% stop_words$word) %>%
438     filter(!is.na(word1) & !is.na(word2)) %>%
439     unite(bigram, word1, word2, sep = " ")
440
441   # Calculate TF-IDF
442   bigrams %>%
443     count(!sym(grouping_var), bigram) %>%
444     bind_tf_idf(bigram, !sym(grouping_var), n) %>%
445     arrange(desc(tf_idf))
446 }
447
448 # Analyze TF-IDF if neighborhood_cleansed is available
449 if ("neighborhood_cleansed" %in% names(airbnb)) {
450   neighborhood_tfidf <- analyze_tfidf(
451     airbnb,
452     "neighborhood_overview",
453     "neighborhood_cleansed"
454   )
455
456   # Display top TF-IDF terms by neighborhood
457   top_tfidf_terms <- neighborhood_tfidf %>%
458     group_by(neighborhood_cleansed) %>%
459     slice_max(order_by = tf_idf, n = 5) %>%
460     ungroup()
461
462   print(top_tfidf_terms)
463 }
464
465 }
```

Appendix 9.J: Correlation Analysis

```
466 #----- CORRELATION ANALYSIS -----
467 #----- CORRELATION ANALYSIS -----
468 #----- CORRELATION ANALYSIS -----
469 # Analyze correlations between numeric variables
470
471 # Function to calculate correlations between numeric variables
472 calculate_correlations <- function(data) {
473   # Select numeric columns
474   numeric_cols <- data %>%
475     select(where(is.numeric)) %>%
476     # Remove columns with too many NAs
477     select_if(function(x) mean(!is.na(x)) > 0.5)
478
479   # Calculate correlation matrix
480   cor_matrix <- cor(numeric_cols, use = "pairwise.complete.obs")
481
482   # Convert to long format for visualization
483   cor_data <- as.data.frame(cor_matrix) %>%
484     rownames_to_column("variable1") %>%
485     pivot_longer(-variable1, names_to = "variable2", values_to = "correlation")
486
487   return(cor_data)
488 }
489
490 # Calculate correlations
491 correlations <- calculate_correlations(airbnb)
492
493 # Visualization: Correlation Heatmap
494 plot_correlation_heatmap <- function(cor_data) {
495   ggplot(cor_data, aes(x = variable1, y = variable2, fill = correlation)) +
496     geom_tile() +
497     scale_fill_gradient2(
498       low = "blue",
499       mid = "white",
500       high = "red",
501       midpoint = 0,
502       limits = c(-1, 1)
503     ) +
504     theme(
505       axis.text.x = element_text(angle = 45, hjust = 1),
506       axis.title = element_blank()
507     ) +
508     labs(
509       title = "Correlation Heatmap of Numeric Variables",
510       fill = "Correlation"
511     )
512 }
513
514 # Generate and display correlation heatmap
515 plot_correlation_heatmap(correlations)
```